

Azure Kubernetes Service

The BIG Picture

Notes based on
Pluralsight course

You will learn

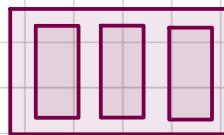
- What is AKS?
- AKS in Action

Other Courses

Azure Container Instances
Designing a compute strategy for Microsoft Azure

1. What is AKS?

You should know about "Containers"



② The case for kubernetes

Monolith	Traditional software Development			Migrate (if hardware fails)
	Build	Package	Deploy	
	Present	Microservices		

Independent, Distributed, etc.

Horizontal vs. Vertical scaling!

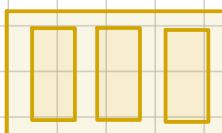
Scale specific features of the service!

Explore the pros and cons of microservices!

"We need automation, which includes automatic scheduling, automatic configuration, supervision and failure handling".



⑥ Containers 101



Containers → lightweight design of a process
isolation of an operating system.

Containers vs. Virtualization

Containers

use namespaces offer isolation of system resources

Filesystems

Process IDs

User IDs

Network Interfaces

} Single Namespace



The process will only see resources that are inside the same namespace.

Process doesn't belong to one namespace but one namespace of each kind.

Namespace kinds

Mount (mnt)

Process ID (pid)

Network (net)

Inter process communication (ipc)

UTS

User ID (user)

Srihari Sridharan

Q3) UTS namespace determines what hostname and domain name the process running inside that namespace sees.

What network namespace a process belongs to determines which network interfaces the application running inside the process sees.

Each container uses its own network namespace and therefore each container sees its own set of network interfaces.

* How can you limit the amount of system resources a container can consume?

Answer: Cgroups → Control Groups

Linux kernel feature that limits the resource usage of a process (or a group of processes)

A process cannot use more than the configured amount of CPU, memory, network bandwidth and so on.

Cgroups prevent processes from hogging resources reserved for other processes.

@asksrhari

DOCKER

Platform for packaging, distributing and running applications.

Host isolation

Simplifies packaging application and its dependencies

Consistent Experience

BENEFITS

Offers Images

Images are much smaller compared to VM
Namespaces and cgroups to offer isolation and resource controls

Images

The image contains the filesystem and the path to the application executable when you run the image.

Registries

A Docker registry is a repository that stores your Docker images and allows easy sharing of those images between different people and computers.

Container

A running container is a process running on the host running Docker, but it's completely isolated from host and other processes.

Srihari Sridharan

Docker Workflow

① Build image &
push to registry

② Pull image ③ Run
image

See

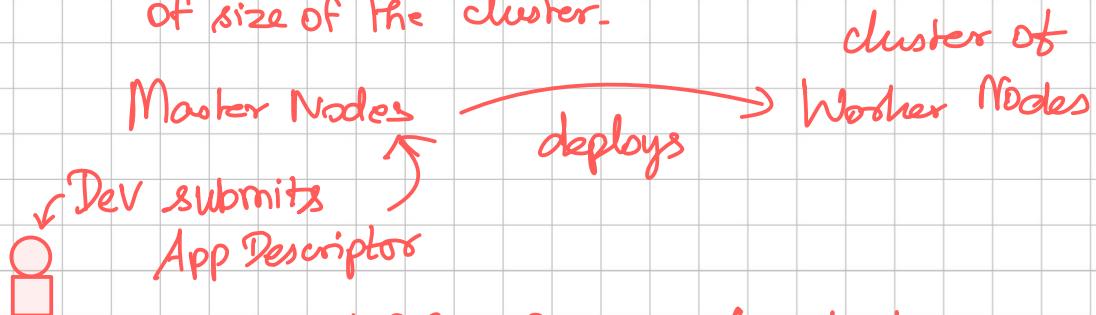
Container Management
using Docker.

C) What is Kubernetes?

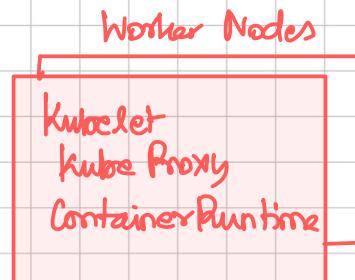
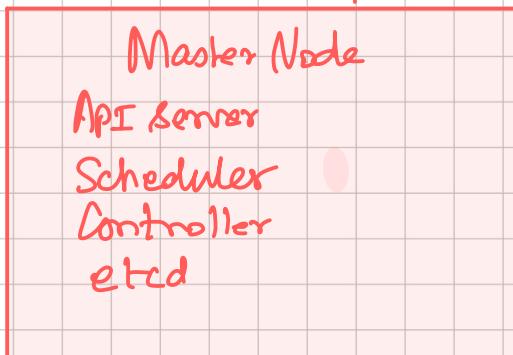
Allows you to easily deploy and manage containerized applications on top of it.

Exposes the underlying infrastructure as a single computational resource.

Consistent deployment experience regardless of size of the cluster.



K8S Cluster Architecture



Running Applications in k8s

1. Package your appn. in to one or more containers
 2. Push those images to an image registry
 3. Post app descriptors to k8S API server
 4. Scheduler schedules the containers on available workers
 5. Kubelet instructs nodes to download container images
 6. Kubelet instructs nodes to run the containers

Benefits

Simplifies appn. deployment

Better hardware utilization

Health monitoring and self healing
Automatic scaling.

Automatic scaling -

Kubernetes Objects

`kubectl <operation> <object> <resource name>`
`<optional flags>`

Pod

smallest unit that k8s manages

Made up of one or more containers

Querying a pool returns a data structure that contains information about containers and its metadata.

Pods aren't durable!

@asksrhari All containers for a pod will be running on the same node
Any container running within a pod will share the Node's network with any other containers in the same pod
Containers within a pod can share files through volumes, attached to containers.

A pod has an explicit lifecycle and will remain in a node it was started.

Namepaces

Pods are collected into namespaces, which are used to group pods.

Namespaces can be used to provide quotas and limits around resource usage and have an impact on DNS names that k8s creates internal to the cluster.

If no namespace is mentioned k8s assumes you are working with default namespace.

Nodes

Node is a machine that is added to the kubernetes cluster.

@asksrihari The master node is the brain of k8s while the worker nodes do the actual work of pulling container images and running pods.

Networks All the containers in a Pod share the Node's network

All nodes in k8s cluster are expected to be connected to each other and share a private cluster-wide network.

k8s runs containers within a pod within this isolated network.

Handling IPs, DNS entries etc.

Controllers

Desired state deployment
Controllers are the brain behind this

ReplicaSets Associated with a Pod and indicates how many instances of that Pod should be running within the cluster
also implies a controller that watches ongoing state and knows how many of your pod to keep running.

Replica set is commonly wrapped in turn by a deployment.

Deployment

Recommended way to run code on k8s



wraps around and extends the replica set

Includes metadata settings to know how many pods to keep running.

Services

k8s resource used to provide an abstraction through to your pod agnostic of the specific instances that are running.
Can contain a policy.

Emulates a software load balancer within Kubernetes.

② What is AKS ?

Self hosting k8s cluster

Manually install the master and worker nodes.

LOT OF WORK !

Handled by

AKS

Need to consider high availability of the master, adding additional worker nodes, patching, upgrades, etc.

Aks. Deployment , management and operations of kubernetes

Provisioning , upgrading and scaling on demand.

Master node is managed by Azure

Offload responsibility to Azure

Pay for agents only

Kubectl is part of cloud shell.

Benefits

- Version upgrades and patching
- Easy cluster scaling
- Self healing hosted control plane (master)
- Cost savings

F) Beyond Managed k8s

Think about the ecosystem of the cloud vendor.

Azure Container Instances (ACI)

Azure Container Registry (ACR)

Azure Service Fabric (SF)

Azure App Service (AAS)

Azure Batch Service (ABS)

A CI

Deploy containers without worrying about underlying infrastructure

Easily start deploying containers for targeted usecases.

- 'Dockerize' your application and execute it in one click.

ACR

Repository for container images.

SF

Foundational technology powering core Azure infrastructure as well as other Microsoft services.

Highly available and durable services at cloud scale.

Microsoft's container orchestrator.

AAS

PaaS Platform

Built in autoscaling and load balancing and CI/CD with GitHub.

ABS

Cloud scale job scheduling and compute management.

Batch pool to run tasks in Docker containers.

Tool Batch Shipyard!

@asksrihari

2. AKS in Action

@ Quick tour of the sample app.

<https://github.com/monojnair/myapp>

Docker images

<https://hub.docker.com/r/monojnair/myapp/tags>

Run

> docker container run --name myapp -d -p 8081:80
monojnair/myapp:v1

visit <http://localhost:8081>

Similarly run V2, V3 and V4 in 8082, 8083 and 8084

NOTE: You can use the ones above or write your own apps and build images

④ Deploying on AKS cluster

RG aks-rg1 Stuff using Azure portal / CLI

Create k8s cluster

Various blocks in portal

Basic | Node Pools | Authentication | Networking

↓ ↓
Set defaults

Better to use

System assigned

Managed identity.

Integration

Tags

Relieves the burden of renewing the

Srihari Sridharan

@asksrhari credentials for service principal.

RBAC : enable

Network → Azure CNI
Configuration

↳ Look at
course

Implementing Managed
Identities for Microsoft
Azure Resources

Integration

Azure monitor default workspace for
log analytics.

Tags Tag your resources as required

In azure CLI

az account show

az configure --defaults group=aks-rg1

Set RG
for subsequent
CLI commands

az aks get-credentials --name <cluster
name>

(eg) az aks get-credentials --name aksdemol

Fetch the
credentials and
merge it into current
context.

Whatever you used
in Basic section.

az aks install-cli ← to install kubectl

③ Deploy the application to AKS cluster imperatively.

K8s management techniques.

Imperative - for development projects how learning curve

Imperative object configuration - for production projects moderate learning curve

Declarative object - for production configuration projects high learning curve

This sample uses (imperative) approach

```
kubectl create deployment myapp  
--image=manojnair/myapp:v1  
--replicas=1
```

Service

Creating a service of type load balancer
In AKS this creates an azure load balancer

@asksrihari

```
kubectl expose deployment myapp --type=LoadBalancer  
--port=80 --target-port=80
```

↑ Before running the above command
have another terminal with watch
enabled.

```
kubectl get svc --watch
```

Creates the load balancer!

④ Scaling the deployment manually

To scale deployments

```
kubectl scale deployment myapp --replicas=3
```

Application is powered by 3 replicas!

To get max pods that can be supported by
a cluster and agent pool

```
az aks nodepool show --cluster-name akademil  
--name agentpool --query "maxPods"
```

Note If maxPods is 110 k8s cannot create
these many pods for your application as

```
kubectl get pods --namespace kube-system
```

Add more worker nodes incase you wish to increase the number of replicas (pods) beyond the node's max pods.

② Scaling nodes manually

`kubectl get nodes` returns the number of nodes.

to scale nodes

```
az aks scale
--resource-group aks-rg1
--name aksdemo1
--nodecount 2
--no-wait
```

returns immediately we don't need to wait!

To check the status

```
az aks nodepool show --name agentpool
--cluster-name aksdemo1
--query "[count, provisioningState]"
```

use

`kubectl get nodes`

f) Updating the application

Rollouts Manage the rollout of a resource
 Valid resource types include deployments,
 daemonsets and statefulsets.

Usage

`kubectl rollout SUBCOMMAND`

To check the rollout history

`kubectl rollout history deployment/myapp`

Describe the deployment

`kubectl describe deployment myapp`

You can find the
 Container and image details

Image manojnair/myapp:v1

Method 1

`kubectl set image`

`deployment /myapp`

`myapp= manojnair /myapp:v2`

`--record = true`

Tracks the change →

Srihari Sridharan

@asksrihari As of 2022 FEB there is a warning that --record will be deprecated.

Refresh the app and you will see v2.

Method 2

```
kubectl edit deployment myapp  
--record=true
```

Refresh the app and you will see v3.

Check `kubectl rollout history deployment/myapp`

(g) Rolling back the applications to previous versions

To undo a deployment

Shows the history of changes.

```
kubectl rollout undo deployment / myapp
```

NOTE: If you keep running this command repeatedly, then you will basically keep toggling between the last 2 deployments.

To revert to specific version in history

```
kubectl rollout undo deployment / myapp  
--to-revision 1
```

IMPORTANT NOTE

When we undo a rollout, K8S creates a new item in history and removes the old entry. (This explains why it keeps swapping between last two rollouts when we undo twice or more)

⑬ Using declarative approach to deploy kubernetes objects

myapp2.yaml

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp2
  labels:
    app: myapp2
spec:
  replicas: 3
  selector:
    matchLabels:
      app: myapp2
  template:
    metadata:
      labels:
        app: myapp2
  spec:
    containers:

```

@asksrihari

```
- name : myapp2
  image: manojnair / myapp : v2
  ports :
    - containerPort : 80
apiVersion : v1
kind : Service
metadata:
  name : myapp2
spec:
  selector:
    app: myapp2
  type: LoadBalancer
  ports:
    - protocol: TCP
      port: 80
      targetPort : 80
```

separator for multiple resources in a file

Pod image

load balancing service

Build k8s objects declaratively and manage using source control

```
kubectl apply -f ./myapp2.yaml
```

To clean up

```
kubectl delete deployment myapp
```

```
kubectl delete service myapp
```

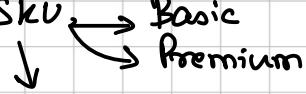
(i) Pushing the images to ACR

Azure Container Registry

ACR + AKS - Best of both worlds!

Create ACR in portal or using IAC.

Registry name, location, SKU



registry-name.azurecr.io

e.g. myappacr01

Networking and encryption are available only for premium SKU.

Validate and create

AKS agent pool nodes need permission to pull images from ACR.

Access Control → Add Role Assignment

Role - reader

Select the system assigned managed identity →

aksdemo1-agentpool identity.

SAVE !

Srihari Sridharan

@asksrihari

Push images to ACR

docker image ls

docker tag manojnair /myapp:v1

myappacr01.azurecr.io/myapp:v1

az acr login --name = myappacr01

docker push myappacr01.azurecr.io /myapp :v1

Create the deployment with the ACR image

kubectl create deployment myappacr

--image = myappacr01.azurecr.io /myapp :v1

This command can be run and the images can be pulled only by agent pool nodes.

kubectl expose deployment myappacr

--type = LoadBalancer

--target-port = 80

--port = 80

Image is securely stored in ACR!

Srihari Sridharan

Next Steps

- Learning paths on Pluralsight k8s
 - Certified Kubernetes Administrator
 - Certified Kubernetes Application Developer
 - Kubernetes Administration
 - Using Kubernetes as a developer
 - Azure Kubernetes Service Workshop
<https://aka.ms/learn/aksworkshop>
-
-

THE
END